

Assignment 2

Siwon Kim

Electrical and Computer Engineering
Seoul National University

<http://data.snu.ac.kr>

Assignment Objectives

- Part 1: Implementing CNN
 - To understand CNN architecture before using the TensorFlow
 - Implement forward / backward passes for (1) convolution layer and (2) max pooling layer
- Part 2: Training CNN
 - Learn how to define, train, and evaluate CNNs with TensorFlow
 - Explore various hyperparameters to design a better CNN model
- Part 3: Visualizing CNN
 - Learn how to visualize and interpret a trained CNN model
 - Implement the codes for generating (1) image-specific class saliency maps, (2) class representative images, and (3) adversarial examples

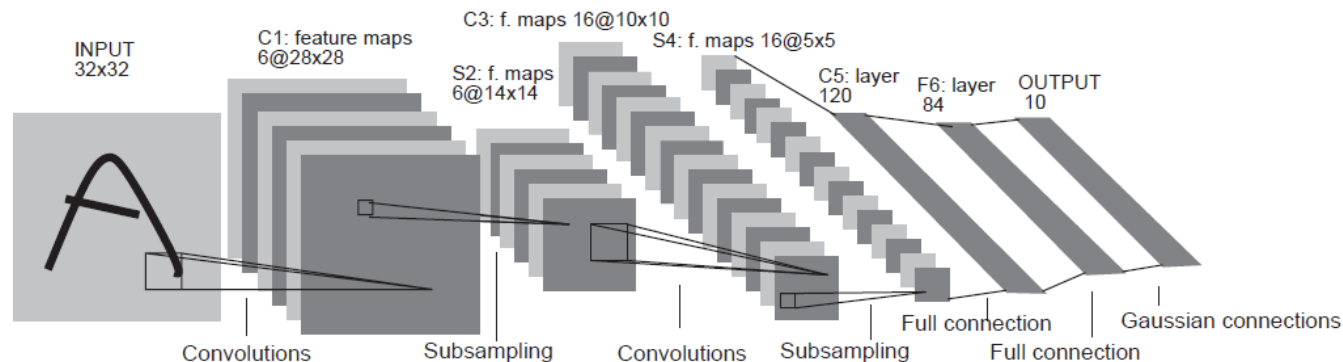
CIFAR10 Dataset




- Consists of $32 \times 32 \times 3$ color images in 10 classes
 - Training set: 40k instances
 - Validation set: 10k instances
 - Test set: 10k instances
- *Model training
 - *Model evaluation and selection
 - *Final model testing

Convolutional Neural Networks (CNNs)

- Simply neural networks that use convolution in their layers



source: [LeCun et al., 2015b]




- Three key ideas behind CNN
 - Local connectivity
 - Invariance to location
 - Invariance to local transition
- Convolution layer
- Pooling layer

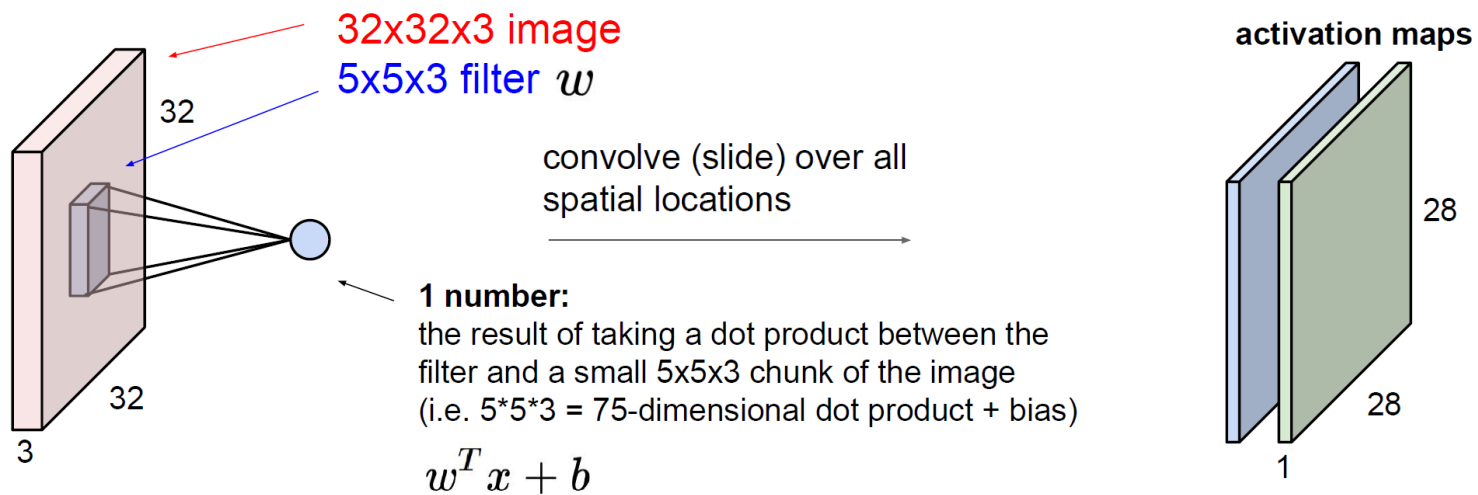
Fully Connected Layer and Convolution Layer

- Fully connected layer

32x32x3 image -> stretch to 3072 x 1




- Convolution layer



Convolution and Cross-correlation


- Convolution

$$\begin{aligned} S(i, j) &= (I * K)(i, j) \\ &= \sum_m \sum_n I(i - m, j - n) K(m, n) \end{aligned}$$



- Cross-correlation

$$\begin{aligned} S(i, j) &= (I * K)(i, j) \\ &= \sum_m \sum_n I(i + m, j + n) K(m, n) \end{aligned}$$



- Many NN libraries implement cross-correlation but call it convolution (without kernel flipping)

Convolution Layer

- Input is convolved with a set of filters, giving feature maps (without kernel flipping)

1 _{x1}	1 _{x0}	1 _{x1}	0	0
0 _{x0}	1 _{x1}	1 _{x0}	1	0
0 _{x1}	0 _{x0}	1 _{x1}	1	1
0	0	1	1	0
0	1	1	0	0

Image

4		

Convolved
Feature



$$\begin{bmatrix} -1 & -1 & -1 \\ -1 & -8 & -1 \\ -1 & -1 & -1 \end{bmatrix}$$




$$\begin{bmatrix} 0 & -1 & 0 \\ -1 & 5 & -1 \\ 0 & -1 & 0 \end{bmatrix}$$



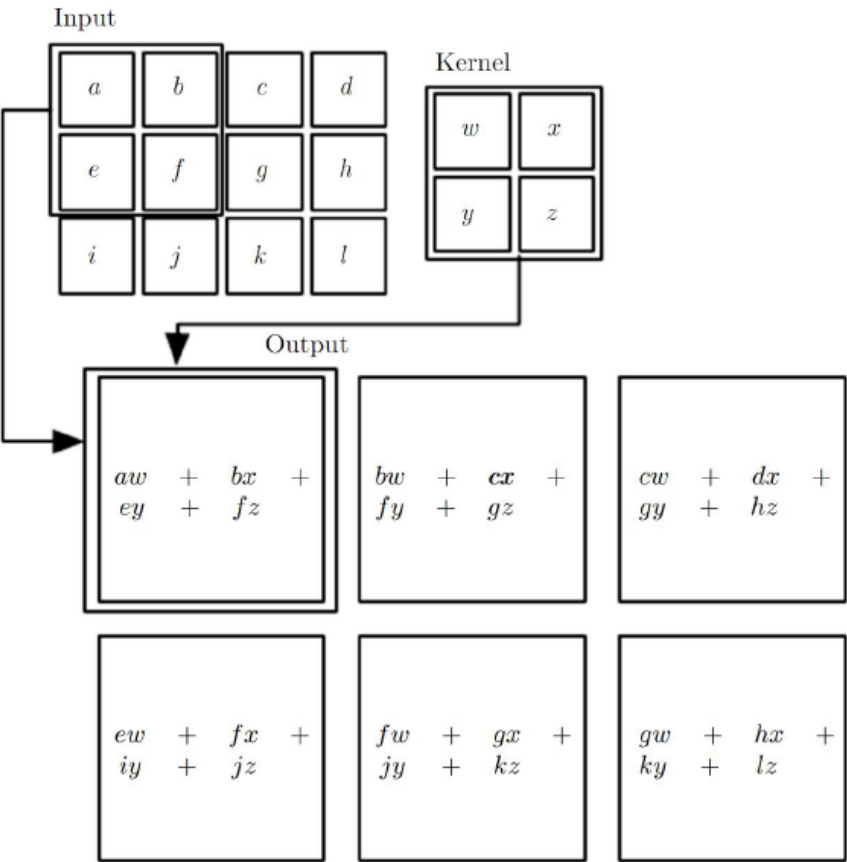
source: www.wikipedia.org

- Filter (aka kernel or convolution matrix)
 - Different filters extract different features (e.g. edges)
 - Filter weights: trained with data
 - Weight sharing & local connectivity

Convolution Example




source: [Angermueller et al., 2016]



source: [Goodfellow et al., 2016]


Convolution with Stride

- The number of pixels between adjacent receptive fields
= down-sampling the output of full convolution function



7x7 input
3x3 filter with stride 1

→ 5x5 output



7x7 input
3x3 filter with stride 2


→ 3x3 output

Zero-padding

- Implicitly zero-pad input to make it wider

type	output	# zeros padded		
		left	right	total
same	m	k even	$\lfloor \frac{k-1}{2} \rfloor$	$\lfloor \frac{k-1}{2} \rfloor + 1$
		k odd	$\frac{k-1}{2}$	$\frac{k-1}{2}$
valid	$m - (k - 1)$	0	0	0

(m : input width; k : kernel width; $s = 1$)




- Examples

$$m = 5, k = 3$$

same




valid




$$m = 5, k = 4$$

same





valid



Pooling Layer


- Nonlinear down-sampling
 - Aggregates statistics of local features (with max or average operation)
 - Reduced variance: provides invariance to local transformations



source: [Angermueller et al., 2016, Thériault et al., 2013]

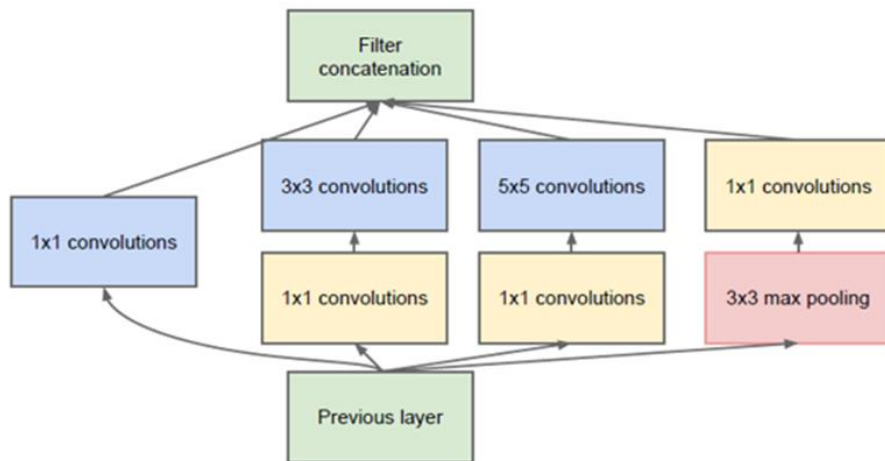
Inception model (a.k.a GoogLeNet)

- Deeper network with computational efficiency
 - ImageNet Large Scale Visual Recognition Competition (ILSVRC) 2014 winner (6.7% top 5 error)
 - 22 layers with 5 million parameters (12x less than AlexNet *ILSVRC 2012 winner)
 - Efficient “Inception” module




Inception module

- Local network topology composing the Inception model
 - Apply parallel filter operations on the input from previous layer
 - Multiple filter sizes for convolution (1x1, 3x3, 5x5)
 - 1x1 convolution for dimensionality reduction



Inception module



1x1 convolution

Image-specific class saliency maps


- How to tell which pixels matter for classification?
- Visualize the degree to which each pixel affects the classification
- Compute gradient of unnormalized class score with respect to image pixels, take absolute value and max over RGB channels

Given image

Class label: Collie



Saliency map



Class representative images

- Generate an image I^* that achieves a high score for the class y


$$I^* = \arg \max_I s_y(I) - R(I)$$

$$R(I) = \lambda \|I\|_2^2$$

- Starting with a random noise, perform gradient ascent on a target class
- **L2 regularization** and **periodic Gaussian blur** regularization

Class representative image

Target class: Collie



Adversarial examples

- Make the CNN model to miss-classify a given image into a target class
- Starting with a given image (cf. random noise), perform gradient ascent over the image to maximize the target class score
- Stop when the network classifies the image as the target class
- **L2 regularization** to normalize the gradients


Given image
Class label: Collie




Adversarial example
Target class: Tarantula



Difference



Magnified difference (10x)



How To Install Assignment Files

- Assignment files
 - Utils/*
 - Assignment2-1_Implementing_CNN.ipynb
 - Assignment2-2_Training_CNN.ipynb
 - Assignment2-3_Visualizing_CNN.ipynb
 - CollectSubmission.sh
- Install assignment files
 - `tar zxvf assignment2.tar.gz`
 - `sudo chmod 755 CollectSubmission.sh`
 - jupyter notebook
- Open the notebooks on your browser and get started

Submitting your work

- Submitting your work
 - DO NOT clear the final outputs
 - After you are done **all three parts**
 - ✓ `$./CollectSubmission.sh team_#`
 - ✓ Upload the `team_#.tar.gz` on ETL
 - ✓ Your `team_#` is in the excel file,
<http://etl.snu.ac.kr/mod/ubboard/article.php?id=724723&bwid=1535905>

Important Notes

- DUE : 10/15/2018, We do not accept late submissions!
- PLEASE read the notes on the notebooks carefully
- Google first before mailing TAs
- TA email: deeplearning.snu@gmail.com

Thank You!

